

# Scalable AnyDSL Molecular Dynamics application with MPI

Scaling the Molecular Dynamics implementation in AnyDSL

Rafael Ravedutti Lucio Machado

Chair of Computer Science 10, Friedrich-Alexander University of Erlangen-Nuremberg

March 19, 2019



## Motivation

Compare AnyDSL scalable application with other state-of-the-art technologies in order to explore benefits in code writing and performance.

# **Outline**

## **AnyDSL**

## **Molecular Dynamics**

## **Proposal**

## **Experimental Results**

# AnyDSL



# AnyDSL

## Framework for development of domain-specific libraries

- Higher-order functions
- Thorin
- Impala
- Partial evaluation

# AnyDSL

```
1  fn main() {  
2      let img = load("dragon.png");  
3      let blurred = gaussian_blur(img);  
4  }
```

# AnyDSL

```
1  fn gaussian_blur(field: Field) -> Field {
2      let stencil: Stencil = { /* ... */ };
3      let mut out: Field    = { /* ... */ };
4
5      for x, y in @iterate(out) {
6          out.data(x, y) = apply_stencil(x, y, field, stencil);
7      }
8
9      out
10 }
```

# AnyDSL

```
1  fn iterate(field: Field, body: fn(int, int) -> ()) -> () {
2      let grid = (field.cols, field.rows, 1);
3      let block = (128, 1, 1);
4
5      with nvvm(grid, block) {
6          let x = nvvm_tid_x() + nvvm_ntid_x() * nvvm_ctaid_x();
7          let y = nvvm_tid_y() + nvvm_ntid_y() * nvvm_ctaid_y();
8          body(x, y);
9      }
10 }
```



# Molecular Dynamics



# Molecular Dynamics

## Pair-wise interaction of particles simulation implemented in AnyDSL

- Cells of particles (bounding boxes)
- Neighborlists
- Cluster of particles
- Target CPU with vectorization instructions and GPU

# Molecular Dynamics

## Steps

1. Initialize grid
2. Initialize clusters
3. Build neighbor lists
4. Compute forces and update particles (for 20 timesteps)
5. Redistribute particles and go back to item 2

# Molecular Dynamics

## Steps

1. Initialize grid
2. Initialize clusters
3. Build neighbor lists
4. Compute forces and update particles (for 20 timesteps)
5. Redistribute particles and go back to item 2

# Molecular Dynamics

## Steps

1. Initialize grid
2. Initialize clusters
3. Build neighbor lists
4. Compute forces and update particles (for 20 timesteps)
5. Redistribute particles and go back to item 2

# Molecular Dynamics

## Steps

1. Initialize grid
2. Initialize clusters
3. Build neighbor lists
4. Compute forces and update particles (for 20 timesteps)
5. Redistribute particles and go back to item 2

# Molecular Dynamics

## Steps

1. Initialize grid
2. Initialize clusters
3. Build neighbor lists
4. Compute forces and update particles (for 20 timesteps)
5. Redistribute particles and go back to item 2

# Proposal





# Proposal

## Goals

- Scalable version of the application
- First: scale application on homogeneous clusters
- In the future: heterogeneous clusters (both CPU and GPU nodes)
- Compare scalable implementation with other state-of-the-art versions

# Proposal

## Steps

- Domain partitioning
- Communication pattern
- Synchronization of cells (every timestep)
- Particle exchange (after redistribution)

# Proposal

## Domain partitioning

- Define configuration of nodes
- Split domain accordingly
- Define current node bounding box
- Include ghost layer cells
- **get\_sync\_timesteps()**

# Domain partitioning

```

1  fn get_node_config(
2      world_size: i32,
3      rank: i32,
4      xcells: i32,
5      ycells: i32,
6      zcells: i32) -> [i32 * 3] {
7
8      let mut gx = 1, gy = 1, gz = 1;
9      let mut min_missing_factor = xcells * ycells * zcells;
10
11     for i in range(1, world_size) {
12         if(world_size % i == 0) {
13             let rem_yz = world_size / i;
14
15             for j in range(1, rem_yz) {
16                 if(rem_yz % j == 0) {
17                     let k = rem_yz / j;
18                     let missing_factor = xcells % i + ycells % j + zcells % k;
19
20                     if(min_missing_factor > missing_factor) {
21                         gx = i;
22                         gy = j;
23                         gz = k;
24                         min_missing_factor = missing_factor;
25                     }
26                 }
27             }
28         }
29     }
30
31     [gx, gy, gz]
32 }

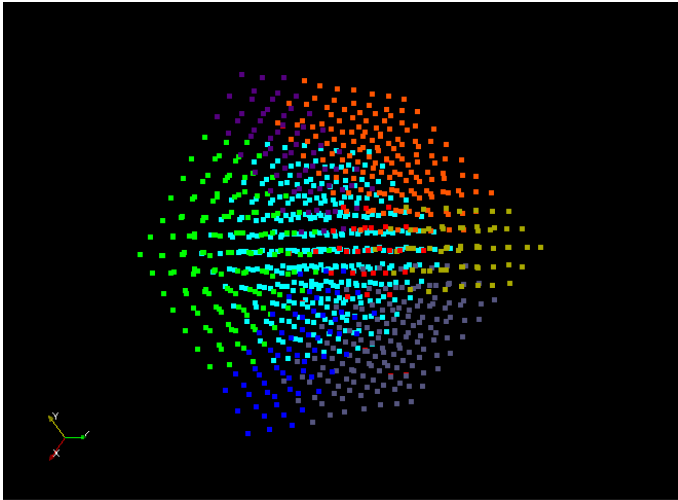
```

## Domain partitioning

```
1  fn @get_node_bounding_box(  
2      world_size: i32,  
3      rank: i32,  
4      cell_spacing: f64,  
5      aabb: AABB) -> AABB {  
6  
7      let mut xmin, xmax, ... : f64;  
8  
9      let xcells = math.floor((aabb.max(0) - aabb.min(0)) / cell_spacing);  
10     let xlength = (xcells / node_dims(0)) * cell_spacing;  
11     let rank_index = unflat_index(rank, ...);  
12  
13     xmin = aabb.min(0) + xlength * rank_index(0);  
14     xmax = aabb.min(0) + xlength * (rank_index(0) + 1);  
15  
16     if(rank_index(0) > 0) {  
17         xmin -= get_sync_timesteps() * cell_spacing;  
18     }  
19  
20     if(rank_index(0) < node_dims(0) - 1) {  
21         xmax += get_sync_timesteps() * cell_spacing;  
22     }  
23  
24     /* y and z are analogous to x */  
25  
26     AABB {  
27         min: [xmin, ymin, zmin],  
28         max: [xmax, ymax, zmax]  
29     }  
30 }
```

# Proposal

## Domain partitioning



# Proposal

## Communication pattern

- Higher-order function for iteration
- Easy to write and change with AnyDSL

# Communication pattern

```
1  fn communication_nodes (
2      world_size: i32,
3      rank: i32,
4      grid: Grid,
5      body: fn(i32, [i32 * 3], [i32 * 3], [i32 * 3], [i32 * 3]) -> () -> () {
6
7      ...
8
9      if(rank_index(0) < node_dims(0)) {
10         @body(
11             flat_index([rank_index(0) + 1, rank_index(1), rank_index(2)], ...),
12             send_begin1, send_end1, recv_begin1, recv_end1 // regions to communicate
13         );
14     }
15
16     if(rank_index(0) > 0) {
17         @body(
18             flat_index([rank_index(0) - 1, rank_index(1), rank_index(2)], ...),
19             send_begin2, send_end2, recv_begin2, recv_end2 // regions to communicate
20         );
21     }
22
23     ... /* y and z are analogous to x */
24 }
```



# Proposal

## Synchronization of cells

- Update positions, velocity and forces of particles
- One-step communication
- Every **get\_sync\_timesteps()** timesteps

## Synchronization of cells

```
1  fn synchronize_ghost_layer_cells(  
2      grid: &mut Grid,  
3      accelerator_grid: AcceleratorGrid,  
4      world_size: i32,  
5      world_rank: i32) -> () {  
6  
7      ... /* Transfer data from accelerator to CPU */  
8  
9      for exchange_rank, send_begin, send_end, recv_begin, recv_end in  
10         communication_nodes(world_size, world_rank, *grid) {  
11  
12         pack_ghost_layer_cells(..., send_begin, send_end);  
13  
14         mpih.irecv(...);  
15         mpih.send(...);  
16         mpih.wait(...);  
17  
18         unpack_ghost_layer_cells(..., recv_begin, recv_end);  
19     }  
20  
21     ... /* Transfer data from CPU to accelerator */  
22 }
```

# Proposal

## Particle exchange

- Exchange redistributed cells
- May be a two-step communication ( $N' > N + N/2$ )
- Every 20 timesteps (redistribution)

## Particle exchange

```
1  fn exchange_ghost_layer_particles(  
2      grid: &mut Grid,  
3      world_size: i32,  
4      world_rank: i32) -> () {  
5  
6      ...  
7  
8      for exchange_rank, send_begin, send_end, recv_begin, recv_end in  
9          communication_nodes(world_size, world_rank, *grid) {  
10         ...  
11  
12         pack_ghost_layer_particles(..., &mut rmng_send_ptcs);  
13         mpih.irecv(...);  
14         mpih.send(...);  
15         mpih.wait(...);  
16         unpack_ghost_layer_particles(..., &mut rmng_recv_ptcs);  
17  
18         if (rmng_recv_ptcs > 0) {  
19             mpih.irecv(...);  
20         }  
21  
22         if (rmng_send_ptcs > 0) {  
23             pack_ghost_layer_particles(...);  
24             mpih.send(...);  
25         }  
26  
27         if (rmng_recv_ptcs > 0) {  
28             mpih.wait(...);  
29             unpack_ghost_layer_particles(...);  
30         }  
31     }  
32 }
```

# Experimental Results



# Experimental Results

## Cluster configuration

- 36-CPU High-Performance-Cluster
- 8 compute nodes
- 4 x Intel(R) Xeon(R) CPU E7-4830, 2.13 GHz 2.4GHz (max. turbo) (8 cores + SMT), SSE 4.1/4.2, 24 MB shared cache
- 256 GB RAM
- 2 x 300 GB SAS internal disks
- NVIDIA GeForce GTX 680
- QDR Infiniband network

# Experimental Results

## Cluster mapping

- MPI processes are mapped to different CPU cores in the same cluster node
- We bind processes to sockets for better memory usage

# Experimental Results

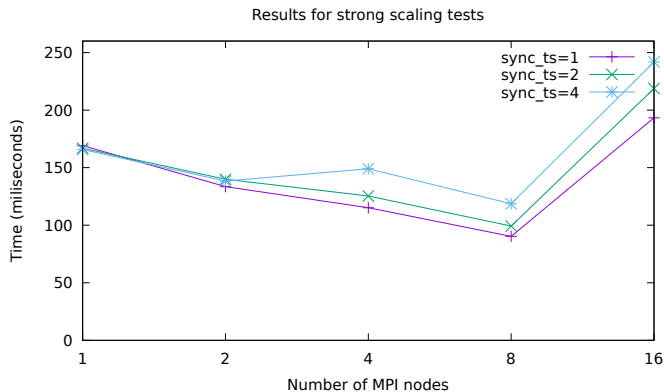
## Strong scaling

- Grid size: 128x128x128
- About 2 million particles




# Proposal

## Strong scaling time



# Proposal

## Strong scaling eff




# Experimental Results

## Weak scaling

- **Problem: cannot use enough particles (too much memory)**
- Must fix this in order to get good weak scaling results


# Experimental Results

## Weak scaling time



# Experimental Results

## Weak scaling efficiency



## Future work

- Reduce memory usage by application
- Improve partitioning for heterogeneous clusters
- Compare with state-of-the-art applications

Thanks for listening.  
**Any questions?**

# References





## References I

- [1] J. Schmitt, H. Kostler, J. Eitzinger, and R. Membarth, “Unified code generation for the parallel computation of pairwise interactions using partial evaluation”, *17th International Symposium on Parallel and Distributed Computing*, 2018.